

Loan Default Prediction

Objective

The objective of the model is to build a classification model which can predict a loan default. We will be interested in the prediction and understanding what are the major factors results to a default.

About Data

The dataset is taken from Kaggle. The dataset consists of multiple deterministic factors like borrower's income, gender, loan purpose etc.

Loading the data.

	ID	year	loan_limit	Gender	approv_in_adv	loan_type	loan_purpose	Credit_Worthiness
0	24890	2019	cf	Sex Not Available	nopre	type1	p1	l1
1	24891	2019	cf	Male	nopre	type2	p1	l1
2	24892	2019	cf	Male	pre	type1	p1	l1
3	24893	2019	cf	Male	nopre	type1	p4	l1
4	24894	2019	cf	Joint	pre	type1	p1	l1

5 rows × 34 columns

(148670, 34)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 148670 entries, 0 to 148669

Data columns (total 34 columns):

#	Column	Non-Null Count	Dtype
0	ID	148670 non-null	int64
1	year	148670 non-null	int64
2	loan_limit	145326 non-null	object
3	Gender	148670 non-null	object
4	approv_in_adv	147762 non-null	object
5	loan_type	148670 non-null	object
6	loan_purpose	148536 non-null	object
7	Credit_Worthiness	148670 non-null	object
8	open_credit	148670 non-null	object
9	business_or_commercial	148670 non-null	object
10	loan_amount	148670 non-null	int64
11	rate_of_interest	112231 non-null	float64
12	Interest_rate_spread	112031 non-null	float64
13	Upfront_charges	109028 non-null	float64
14	term	148629 non-null	float64
15	Neg_ammortization	148549 non-null	object
16	interest_only	148670 non-null	object
17	lump_sum_payment	148670 non-null	object
18	property_value	133572 non-null	float64
19	construction_type	148670 non-null	object
20	occupancy_type	148670 non-null	object
21	Secured_by	148670 non-null	object
22	total_units	148670 non-null	object
23	income	139520 non-null	float64
24	credit_type	148670 non-null	object
25	Credit_Score	148670 non-null	int64
26	co-applicant_credit_type	148670 non-null	object
27	age	148470 non-null	object
28	submission_of_application	148470 non-null	object
29	LTV	133572 non-null	float64
30	Region	148670 non-null	object
31	Security_Type	148670 non-null	object
32	Status	148670 non-null	int64
33	dtir1	124549 non-null	float64

dtypes: float64(8), int64(5), object(21)

memory usage: 38.6+ MB

	ID	year	loan_amount	rate_of_interest	Interest_rate_spread	Upfront_c
count	148670.000000	148670.0	1.486700e+05	112231.000000	112031.000000	109028.0
mean	99224.500000	2019.0	3.311177e+05	4.045476	0.441656	3224.
std	42917.476598	0.0	1.839093e+05	0.561391	0.513043	3251
min	24890.000000	2019.0	1.650000e+04	0.000000	-3.638000	0.0
25%	62057.250000	2019.0	1.965000e+05	3.625000	0.076000	581.4
50%	99224.500000	2019.0	2.965000e+05	3.990000	0.390400	2596.4
75%	136391.750000	2019.0	4.365000e+05	4.375000	0.775400	4812.5
max	173559.000000	2019.0	3.576500e+06	8.000000	3.357000	60000.0

	loan_limit	Gender	approv_in_adv	loan_type	loan_purpose	Credit_Worthiness	open_cr
count	145326	148670	147762	148670	148536	148670	148
unique	2	4	2	3	4	2	
top	cf	Male	nopre	type1	p3	l1	r
freq	135348	42346	124621	113173	55934	142344	148

4 rows x 21 columns

ID	148670
year	1
loan_limit	2
Gender	4
approv_in_adv	2
loan_type	3
loan_purpose	4
Credit_Worthiness	2
open_credit	2
business_or_commercial	2
loan_amount	211
rate_of_interest	131
Interest_rate_spread	22516
Upfront_charges	58271
term	26
Neg_ammortization	2
interest_only	2
lump_sum_payment	2
property_value	385
construction_type	2
occupancy_type	3
Secured_by	2
total_units	4
income	1001
credit_type	4
Credit_Score	401
co-applicant_credit_type	2
age	7
submission_of_application	2
LTV	8484
Region	4
Security_Type	2
Status	2
dtir1	57
dtype: int64	

ID	0
year	0
loan_limit	3344
Gender	0
approv_in_adv	908
loan_type	0
loan_purpose	134
Credit_Worthiness	0
open_credit	0
business_or_commercial	0
loan_amount	0
rate_of_interest	36439
Interest_rate_spread	36639
Upfront_charges	39642
term	41
Neg_ammortization	121
interest_only	0
lump_sum_payment	0
property_value	15098
construction_type	0
occupancy_type	0
Secured_by	0
total_units	0
income	9150
credit_type	0
Credit_Score	0
co-applicant_credit_type	0
age	200
submission_of_application	200
LTV	15098
Region	0
Security_Type	0
Status	0
dtir1	24121
dtype: int64	


Observations

- Data contains 148670 rows, 34 columns.
- Out of 34, 21 are objects and 13 are numeric.
- 14 columns have missing values.
- We can ignore `ID` and `year` columns. As `ID` is unique to each entry and all the data collected in a single year (2019).
- `Status` will be our target column.
- We will ignore some of features in the cleaning process.

Target Distribution.

Now let's look at how target variable `Status` is distributed.

```
Status
0    112031
1     36639
Name: count, dtype: int64
```



The Default percentage is: 24.645%

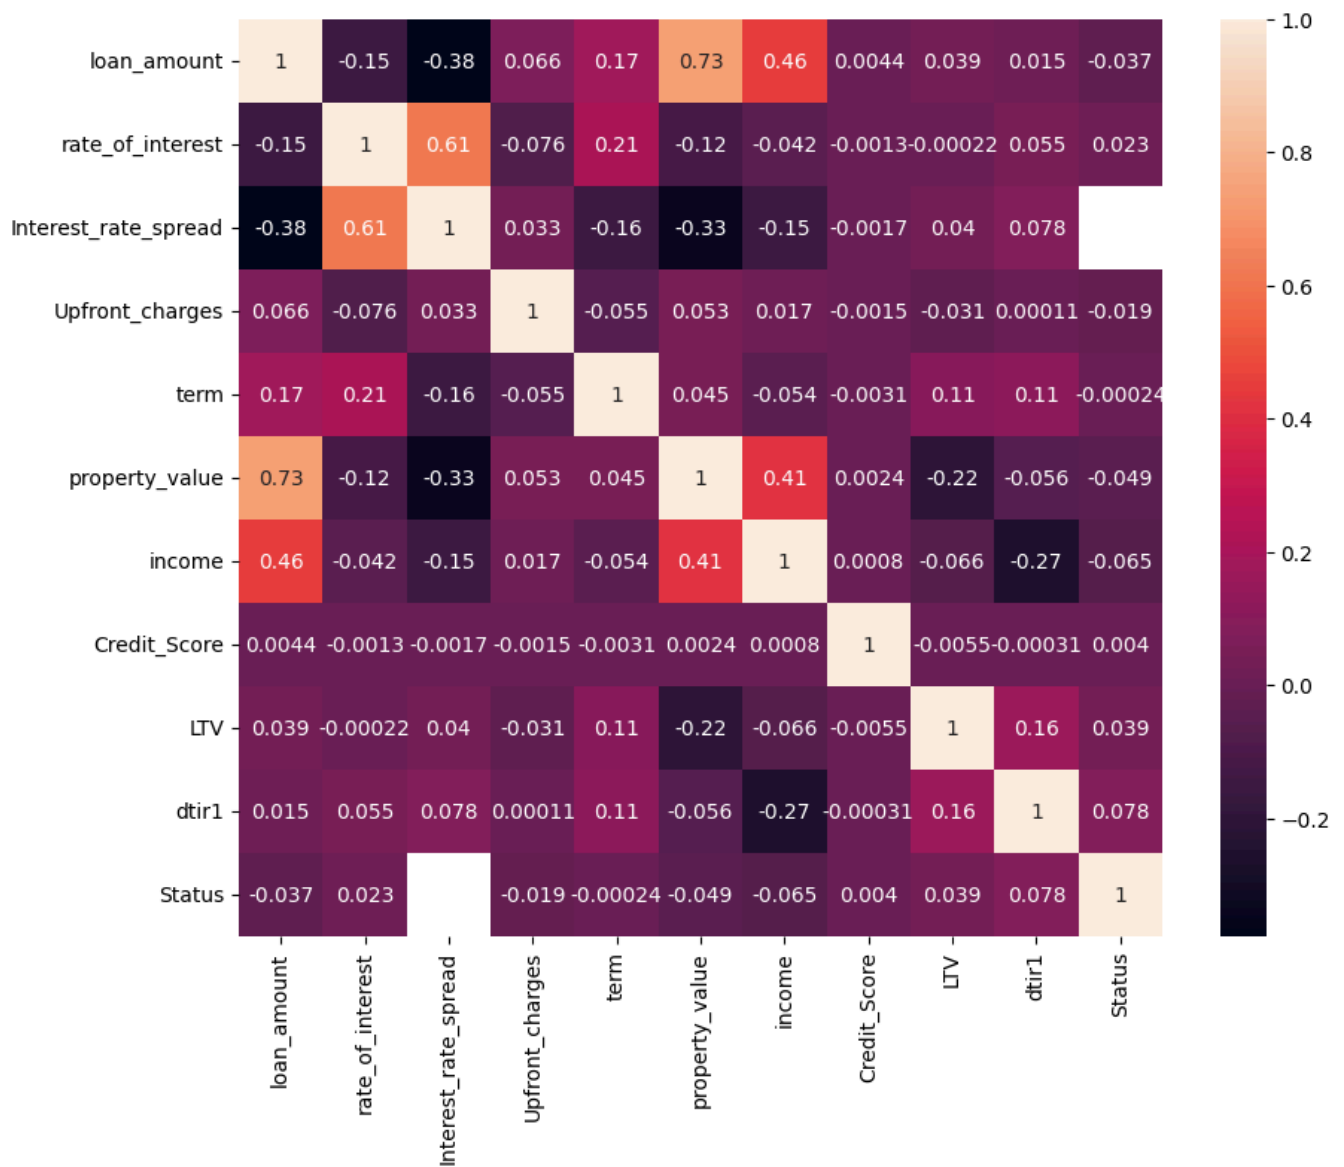
Observations

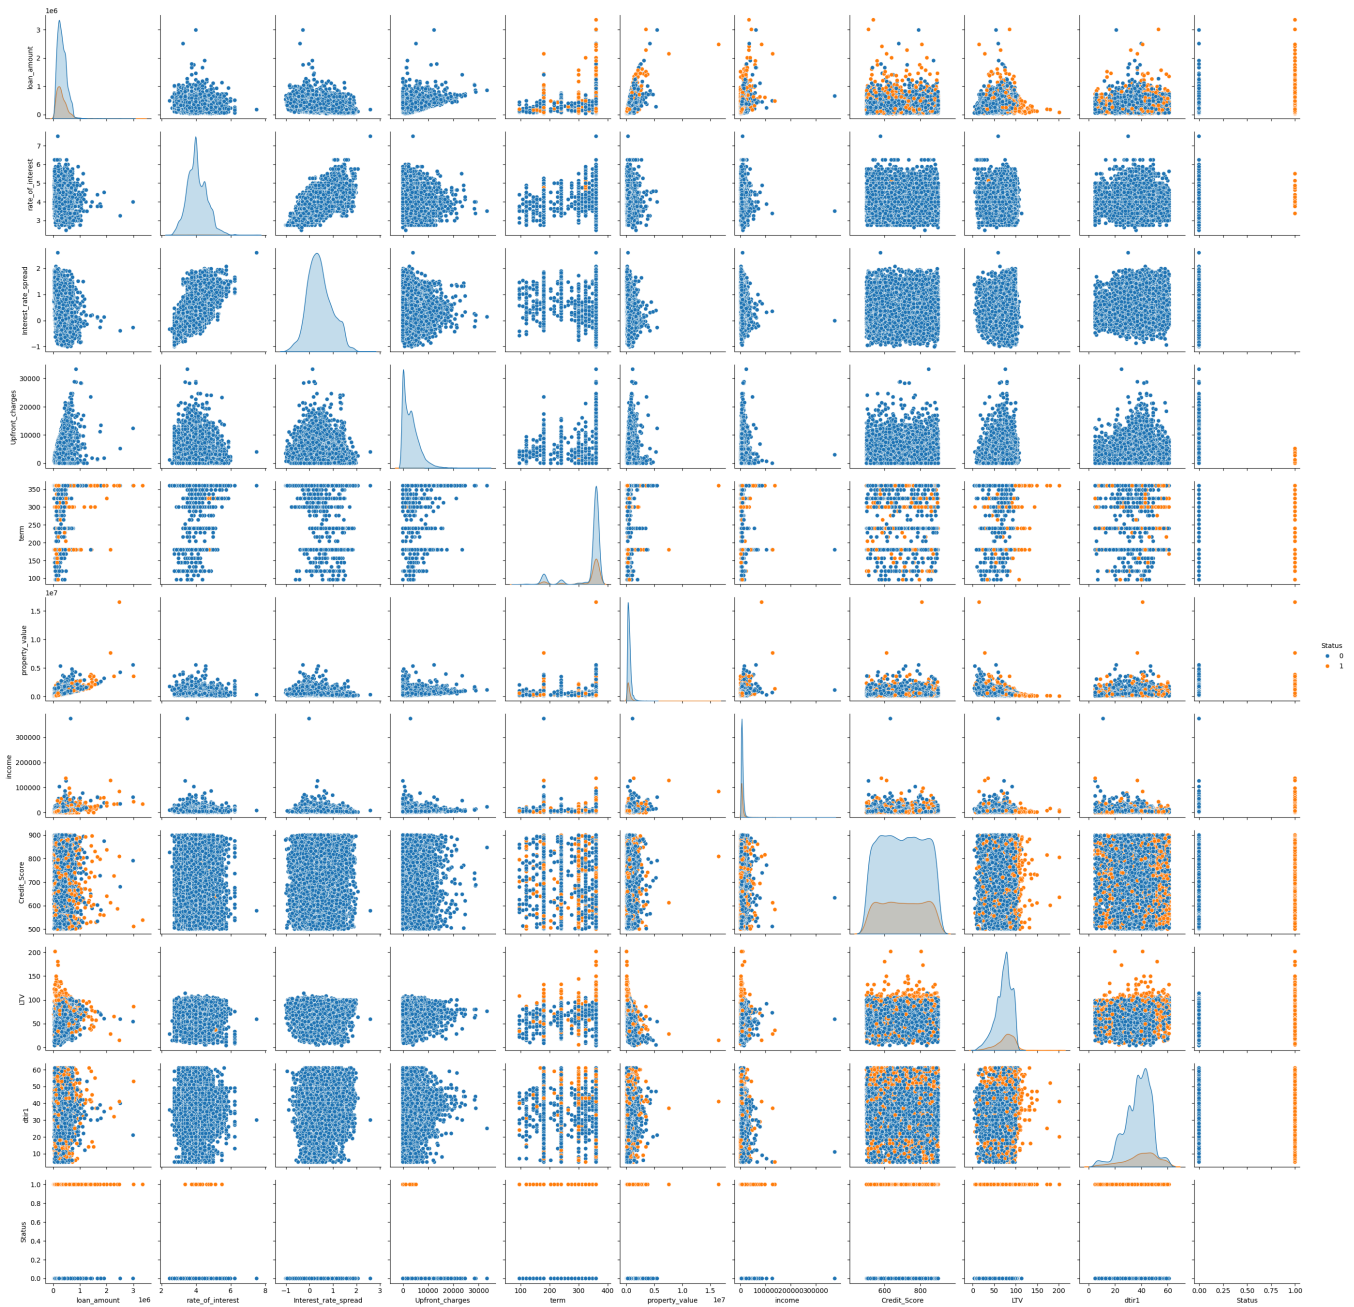
- There are about 25% defaults. It will be helpful if a loan default is predicted early.
- The target data is imbalanced which might affect the model predictions.

Numerical features

	loan_amount	rate_of_interest	Interest_rate_spread	Upfront_charges	term	property_value
0	116500	NaN	NaN	NaN	360.0	118000.0
1	206500	NaN	NaN	NaN	360.0	NaN
2	406500	4.56	0.2000	595.0	360.0	508000.0
3	456500	4.25	0.6810	NaN	360.0	658000.0
4	696500	4.00	0.3042	0.0	360.0	758000.0

ID and year are not considered.





```
Status          1.000000
dtir1           0.078083
income          0.065119
property_value  0.048864
LTV             0.038895
loan_amount     0.036825
Name: Status, dtype: float64
```

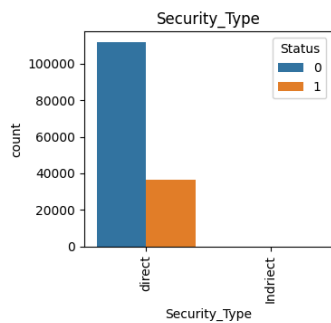
Observations

- We see very little correlation of the features with the targets.
- The top five features correlated to target are
 - `dtir1`
 - `income`

- `property_value`
- `LTV`
- `loan_amount`
- `rate_of_interest` and `interest_rate_spread` are correlated.
- `loan_amount`, `property_value` and `income` are correlated.
- From pair plot we see a almost linear relations between target and features.
- Here we select only the above numerical features to train the model.

Categorical features

Relationship Between Categorical Variables and Target.



- We can ignore the features: `construction_type`, `Security_Type`, `open_credit`, `total_units`, `Secured_by` as they have an even categorical distributions.
- `submission_of_application`, `age` can also be ignored, since we can argue that they are not much influencing features for default.
- `Gender` is also not included.

Selected features

```
num_features = ['income', 'loan_amount', 'dtir1', 'LTV', 'property_value']
cat_features = ['loan_limit', 'approv_in_adv', 'loan_type', 'loan_purpose',
               'Credit_Worthiness', 'business_or_commercial',
               'Neg_ammortization',
               'interest_only', 'lump_sum_payment', 'occupancy_type',
               'credit_type',
               'co-applicant_credit_type', 'Region']
```

Handling Missing Values

Shape of the data: (148670, 34)

Missing values in Numerical features in percentage:

```
dtir1          16.224524
LTV            10.155378
property_value 10.155378
income         6.154571
loan_amount    0.000000
dtype: float64
```

Missing values in Categorical features in percentage:

loan_limit	2.249277
approv_in_adv	0.610749
loan_purpose	0.090133
Neg_ammortization	0.081388
loan_type	0.000000
Credit_Worthiness	0.000000
business_or_commercial	0.000000
interest_only	0.000000
lump_sum_payment	0.000000
occupancy_type	0.000000
credit_type	0.000000
co-applicant_credit_type	0.000000
Region	0.000000

dtype: float64

- There are about 16% missing data in `dtir1` and about 10% in both `LTV` , `property_value` and 6% in `income` .
- `loan_limit` has about 2.25% missing data and less than 1% data is missing in `approv_in_adv` , `loan_purpose` , `Neg_ammortization` .
- We will impute the numerical data with median values categorical data with mode or most frequent values.
- For Numerical data we standardize with standard scalar.
- We one-hot encode the categorical data.
- We create pipelines for these.

Data Preparation

```
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
```

Importing the above from sklearn library.

Creating two transformers (pipelines) each one for categorical and numerical columns.

```
cat_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('one_hot', OneHotEncoder(handle_unknown='ignore'))])
```

```
num_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])
```

Combine these into a column transformer

```
final_transformer = ColumnTransformer(transformers=[
    ('num', num_transformer, num_features),
    ('cat', cat_transformer, cat_features)])
```

Now, another pipeline with the `final_transformer` and the desired model like `LogisticRegression` can be easily constructed, trained and evaluated.

Logistic Regression

Logistic regression with default parameters.

Training the logistic regression model with default parameters.


Accuracy: 0.8611354005515571

Confusion Matrix :
[[22156 219]
 [3910 3449]]

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.99	0.91	22375
1	0.94	0.47	0.63	7359
accuracy			0.86	29734
macro avg	0.90	0.73	0.77	29734
weighted avg	0.87	0.86	0.84	29734

ROC-AUC Score: 0.8257302903524814

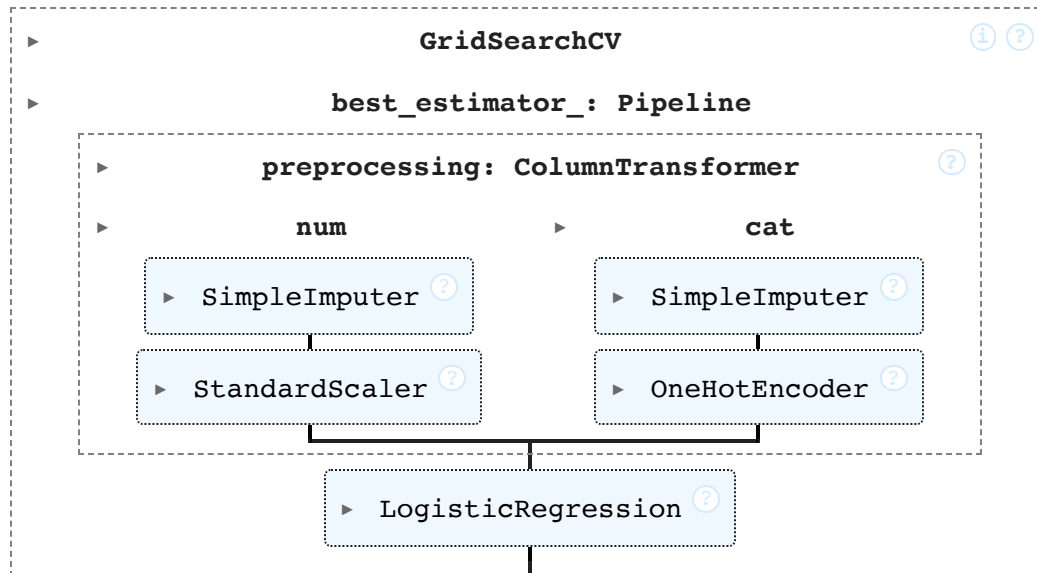


Logistic regression: Finding optimal parameters with Crossvalidation.

Here, GridSearchCV is used to find optimal values for

- Regularization parameter, C
- penalty
- max_iter

f1_score is used for scoring.



Best Parameters: {'classifier_C': 100, 'classifier_max_iter': 100, 'classifier_penalty': 'l2', 'classifier_solver': 'lbfgs'}

Best f1 Score: 0.6315075253200881

Accuracy: 0.861236295150333


Confusion Matrix:

```
[[22155  220]
 [ 3906 3453]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.99	0.91	22375
1	0.94	0.47	0.63	7359
accuracy			0.86	29734
macro avg	0.90	0.73	0.77	29734
weighted avg	0.87	0.86	0.84	29734

ROC-AUC Score: 0.8259000668812029



Observations.

- The best estimator with the tuned parameters has
 - Regularization strength 100
 - Maximum iteration 100
 - penalty 'l2'
 - and solver 'lbfgs'
- The best estimator from the GridSearchCV did not perform any different way than the model with default parameters.
- This model has high precision and low recall. That is when the model predicts the loan will default, it is about 94% of the times correct. But also from low recall of 47%, the model captures about 47% of the actual defaults.

Random Forests

Now, I train a Random forest model and compare the model performance with the Logistic regression. We also tune the hyperparameters using GridSearchCV.

Accuracy: 0.8698795991121275


Confusion Matrix :

```
[[21024 1351]
 [ 2518 4841]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.94	0.92	22375
1	0.78	0.66	0.71	7359
accuracy			0.87	29734
macro avg	0.84	0.80	0.82	29734
weighted avg	0.87	0.87	0.87	29734

ROC-AUC Score: 0.8257302903524814



Observations.

- The Random forest model also has the same accuracy of about 87% almost same as the Logistic regression. I optimize the hyperparameter below.

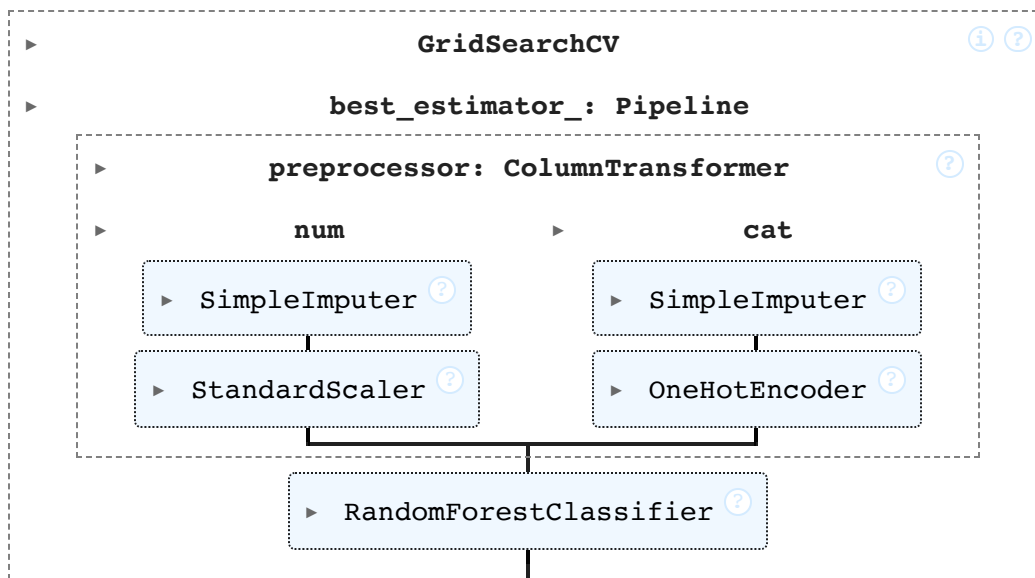
Hyper parameter tuning

The hyperparameters considered for tuning are

- n_estimators
- max_depth
- min_samples_split
- min_samples_leaf

- max_features

Again, f1-score is used for scoring.



Best Parameters: {'classifier__max_depth': None, 'classifier__max_features': 'sqrt', 'classifier__min_samples_leaf': 1, 'classifier__min_samples_split': 10, 'classifier__n_estimators': 300}

Best f1 Score: 0.7375981166982257

Accuracy: 0.8868635232393892


Confusion Matrix:

```
[[21536  839]
 [ 2525 4834]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.90	0.96	0.93	22375
1	0.85	0.66	0.74	7359
accuracy			0.89	29734
macro avg	0.87	0.81	0.83	29734
weighted avg	0.88	0.89	0.88	29734

ROC-AUC Score: 0.881512098209846



Observations

- The best estimator here has
- max_depth: None,
- max_features selected by 'sqrt' method
- min_samples_leaf: 1
- min_samples_split: 10
- n_estimators: 300

Best f1 Score: 0.7375981166982257

- The Random forest has lower precision than the Logistic regression but has better recall and f1-scores.
- The ROC-AUC score is similar to the Logistic Regression.
- We can choose this model, but for only little improvement we sacrifice the interpretability of the model compared to Logistic regression.

Adaboost Classifier

Here, as the third model, I consider Adaboost with Decision tree as the base estimator.

Accuracy: 0.8916055693818524

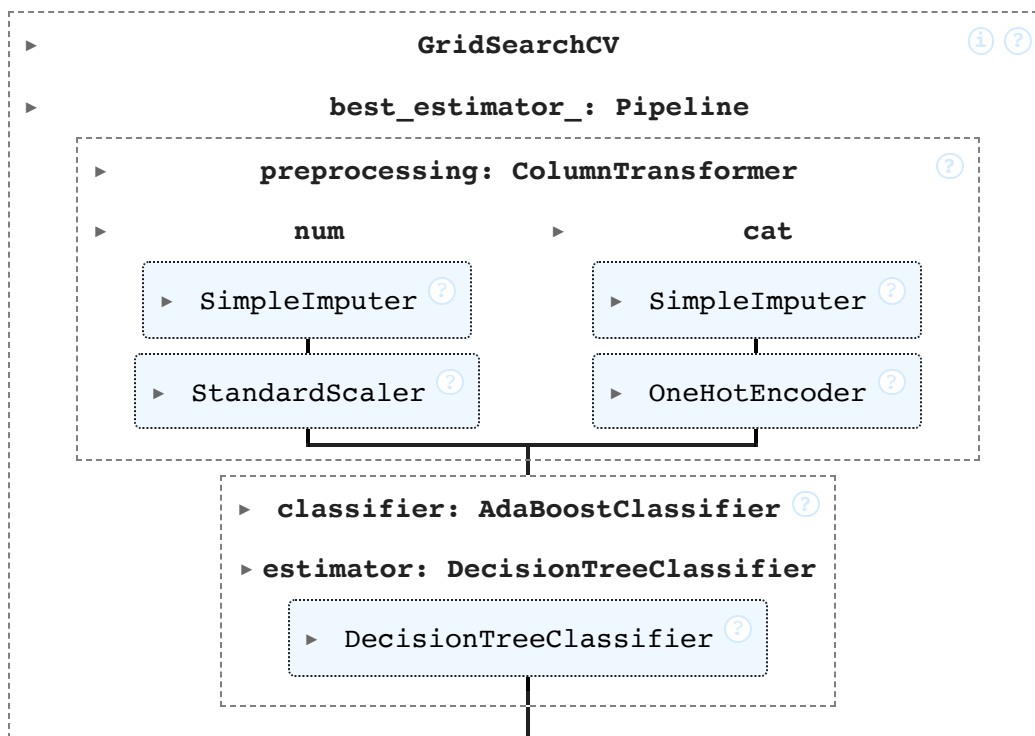
Confusion Matrix :

```
[[21979  396]
 [ 2827 4532]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.98	0.93	22375
1	0.92	0.62	0.74	7359
accuracy			0.89	29734
macro avg	0.90	0.80	0.83	29734
weighted avg	0.89	0.89	0.88	29734

ROC-AUC Score: 0.8801641709577677



Best Parameters: {'classifier__estimator__max_depth': 4, 'classifier__learning_rate': 1.0, 'classifier__n_estimators': 19}

Best f1 Score: 0.7366900250011147

Accuracy: 0.8889823098136813

Confusion Matrix:

```
[[21961  414]
 [ 2887 4472]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.98	0.93	22375
1	0.92	0.61	0.73	7359
accuracy			0.89	29734
macro avg	0.90	0.79	0.83	29734
weighted avg	0.89	0.89	0.88	29734

ROC-AUC Score: 0.8753046844930502

The Best Parameters:

- The base estimator decision tree's max_depth': 4,
 - learning_rate': 1.0,
 - n_estimators': 19 Best f1 Score: 0.7366900250011147
-
- Adaboost improves over Random forest accuracy and precision.
 - This model has the highest f1 and ROC-AUC scores with better predictability.

Which model do I choose?

Here, the problem considers loan defaults. Here model interpretability is very important. Given the scores of all the three models trained here, I would choose the simplest logistic regression model which did almost (may be little less) as good as other two. This is because, it is simple, easy to understand and moreover the model is linear and interpreting the model is very easy.

As a bank, they would like to know what factors affect the most for a given customer to default.

Further steps to improve model

- I can include more features.
- Since the target has unbalanced classes, we can do oversampling or undersampling.
- We can use a different model like SVM, which may perform better.